

The Persona Paradox: Forensic Audit of Persona-Driven Bias and Logical Fragility in LLM-Based Tax Compliance Systems

Ai22btech11027, Ai22btech11013, Ai22btech11006

Abstract—Large language models (LLMs) deployed in high-stakes advisory contexts are increasingly paired with retrieved rule sets and persona-engineered system prompts. This paper investigates whether persona injection produces systematic bias that overrides injected rules—a phenomenon we term *persona dominance*—and quantifies the degree to which persona-conditioned agents diverge from human expert judgment. We construct a controlled experimental framework of 70 tax compliance scenarios across 10 expense kernels and 7 manipulation types, evaluated by 4 persona-conditioned agents under 4 ablation conditions, with human expert ground-truth scores provided for all 70 scenarios. Against this full baseline, the Achievement agent exhibits a strong tendency to over-approve claims that human experts rate as rejectable, while the Security agent over-rejects clearly legitimate claims. Using Fleiss’ Kappa, we show that the full persona-plus-rules system produces inter-agent agreement worse than random chance, while removing personas entirely restores moderate agreement. Decision consistency analysis further suggests that persona prompts exert a stronger effect on outputs than the removal of rules: holding persona constant while removing rules preserves most decisions, whereas holding rules constant while removing personas substantially changes behavior. On contradiction-deliberate scenarios, Achievement consistently assigns high scores despite uniformly low human ground truth, showing a pronounced failure to register disqualifying information. These findings provide empirical evidence that RAG-based rule injection does not reliably function as a behavioral guardrail when persona prompts are present.

Index Terms—Large Language Models, Persona Engineering, Tax Compliance, Retrieval-Augmented Generation, Inter-Rater Reliability, Bias Audit

I. INTRODUCTION

The deployment of LLMs in professional advisory contexts—tax compliance, legal review, financial auditing—proceeds on the assumption that such systems can be constrained through two mechanisms: injection of domain-specific rules via retrieval-augmented generation (RAG), and calibration of evaluative posture through system-prompt persona engineering. A “strict compliance officer” persona paired with a relevant rule set is presumed to produce reliably conservative outputs; an “optimiser” persona paired with the same rules is presumed to produce reliably permissive outputs. The underlying assumption is that rules govern the decision and persona governs the *style* of reasoning about that decision.

This paper challenges that assumption directly. We present evidence that in the presence of a persona prompt, rules function not as decision drivers but as post-hoc justificatory instruments—cited selectively to rationalise decisions effectively

predetermined by the persona’s value orientation. We term this *persona dominance* and demonstrate it through a controlled ablation experiment anchored by human expert ground-truth scores across all 70 scenarios and all 7 manipulation types.

Our experimental context is Indian income tax compliance: evaluation of freelancer and self-employed expense claims across ten expense categories (home office rent, research materials, apparel, travel documentation, hardware assets, subsistence expenses, family payroll, referral fees, vehicle use, and professional development). This domain offers a realistic setting in which rule sets are well-defined, claim legitimacy exists on a spectrum from clearly valid to clearly fraudulent, and the stakes of systematic miscalibration are consequential.

Contributions. First, we introduce a seven-type scenario taxonomy and demonstrate that different manipulation types expose distinct failure modes. Second, we provide formal inter-rater reliability analysis of persona-conditioned LLM agents using Fleiss’ Kappa, showing below-chance agreement ($\kappa = -0.156$) in the full system. Third, we anchor all agent behaviour against a complete human expert baseline covering all 70 scenarios, enabling precise calibration error analysis across every manipulation type. Fourth, we identify a paradoxical RAG effect in which rule injection harms the evaluation of legitimate claims while improving it for illegitimate ones. Fifth, we document Achievement’s complete failure on contradiction and escalation scenarios: in 10 of 10 Contradiction-Deliberate cases, it assigns scores of 4 or 5 against a human GT of 1.

II. RELATED WORK

Prior work on LLM alignment has documented sensitivity to prompt framing [1], persona-induced behavioural shifts [2], and sycophancy under social pressure [3]. Studies of LLMs in legal and compliance contexts have noted tendencies toward overconfident approval of ambiguous claims [4] and inconsistent rule application across semantically equivalent phrasings [5].

Our work extends this literature in three ways. First, we employ Fleiss’ Kappa as a formal reliability metric rather than accuracy measures requiring unambiguous ground truth—appropriate for tax compliance where expert disagreement is common. Second, we isolate the causal contribution of persona versus rules through a four-condition ablation design. Third, we provide a complete human baseline covering all 70

scenarios including contradiction and escalation manipulation types not previously studied in compliance-adjacent LLM evaluation work, enabling calibrated error analysis that distinguishes between failure on legitimate claims (over-rejection) and failure on illegitimate claims (over-approval).

III. EXPERIMENTAL DESIGN

A. Structured Rule Space Design

The knowledge base is constructed as a compact and interpretable rule set derived from common tax principles relevant to small business and freelance scenarios. Rather than attempting to capture the full complexity of tax law, the rule set is deliberately constrained and organized into five thematic buckets: general deductibility principles, mixed-use assets, professional growth expenses, capital versus revenue distinctions, and evidence and compliance requirements.

This organization is functional rather than purely taxonomic. Each bucket introduces a specific type of ambiguity that is later exploited during evaluation. General rules provide baseline interpretive flexibility, mixed-use rules introduce proportional reasoning, professional growth rules test causal linkage, capital versus revenue rules introduce structural classification, and evidence-related rules introduce uncertainty and risk sensitivity. Together, these dimensions ensure that scenarios require multi-step reasoning rather than direct rule lookup.

Each rule is represented as a structured unit consisting of a rule identifier, rule text, explanation, category, and an illustrative example. The inclusion of examples anchors abstract principles in concrete usage patterns, reducing ambiguity and providing implicit templates for application during reasoning.

B. Rule Representation and Construction Process

The rule dataset is constructed as a controlled and interpretable decision space rather than a direct extraction of formal tax law. Each rule follows a structured format consisting of a rule identifier, rule text, explanation, category, and an illustrative example. This representation separates abstraction from application: the rule text captures the principle, the explanation provides contextual interpretation, and the example demonstrates how the rule applies in practice.

The rules are designed around five functional buckets, each introducing a distinct dimension of ambiguity. General rules establish baseline constraints such as business purpose and reasonableness. Mixed-use rules introduce proportional allocation problems. Professional growth rules test whether an expense has a direct linkage to business activity. Capital versus revenue rules introduce classification-based reasoning. Evidence and risk rules impose constraints related to documentation, audits, and compliance.

This construction is deliberate. The objective is not to create independent rules, but to create overlapping constraints that interact. Real-world tax decisions rarely depend on a single rule; instead, they require reconciling multiple conditions simultaneously. The rule set is therefore designed such that a single scenario activates rules from multiple buckets, forcing the model to resolve competing considerations.

The rules are also balanced between permissive and restrictive formulations. Some rules enable claims under certain conditions, while others impose limitations or risks. This balance ensures that the model is not guided toward a single deterministic outcome, but instead must navigate trade-offs between competing objectives such as optimization, compliance, and fairness.

The rules are loosely aligned with principles such as Section 37(1) of the Income Tax Act (business purpose), GST compliance requirements, and general audit and documentation norms. These principles were informed by secondary explanatory resources on Indian small business taxation and cross-referenced with official guidance from the Income Tax Department of India. Rather than reproducing formal legal text, the rules are intentionally simplified and structured to be interpretable and composable across scenarios.

The rule dataset is stored in a structured JSON format to ensure consistency and ease of processing. Each rule is represented as a dictionary with fixed fields, enabling both programmatic retrieval and interpretability. A simplified schema of the dataset is as follows:

```
{
  "rule_id": "RXX",
  "rule_text": "...",
  "explanation": "...",
  "category": "...",
  "example": "...",
  "embedding_text": "..."
}
```

Here, `rule_text` captures the core principle, `explanation` provides contextual clarification, `category` links the rule to its functional bucket, and `example` anchors the rule in a concrete scenario. The `embedding_text` field combines rule text and explanation and is used for semantic embedding generation.

An example rule from the dataset is shown below:

```
{
  "rule_id": "R02",
  "rule_text": "Mixed-use expenses should be claimed proportionally based on business purpose.",
  "explanation": "If an expense or asset is used for both personal and business purposes, only the business portion is eligible.",
  "category": "Mixed Use",
  "example": "A freelancer uses a laptop 60% for work and 40% for gaming."
}
```

This structure ensures that each rule is both machine-readable and semantically rich, supporting accurate retrieval as well as consistent downstream reasoning.

For retrieval, each rule is converted into an embedding input by combining the rule text and explanation. This ensures that semantic similarity reflects both conceptual meaning and contextual interpretation. The example field is excluded from

embedding and instead provided during reasoning, where it functions as an application guide rather than a retrieval signal.

The resulting rule space is compact yet expressive. It remains fully interpretable while being sufficiently rich to produce scenarios involving interacting constraints. This ensures that variability in outputs arises from differences in reasoning rather than inconsistencies in knowledge.

C. Semantic Representation and Embedding

Each rule is converted into a dense vector representation using the Sentence-BERT model `all-MiniLM-L6-v2`. The model is used strictly as a semantic encoder, mapping text into a vector space where similarity reflects meaning rather than surface form.

The embedding input combines rule text and explanation, ensuring that both principle and context influence representation. This allows queries with different phrasing but similar intent to retrieve the same rules.

All embeddings are generated once and stored as fixed representations. This precomputation ensures stability of the semantic space and eliminates variability due to repeated encoding. The embedding matrix therefore acts as a static representation of the knowledge space.

D. Deterministic Retrieval via FAISS

All rule embeddings are indexed using FAISS with an exact inner-product index (`IndexFlatIP`) over L2-normalized vectors, effectively implementing cosine similarity.

The use of exact search is a deliberate design choice. While approximate methods exist, they introduce variability in retrieval. Given the small size of the rule set, exact search is computationally efficient and ensures complete determinism.

At inference time, a scenario is encoded and projected into the same embedding space. The FAISS index retrieves the top- k most similar rules, along with their explanations and examples, implementing a deterministic mapping $S \rightarrow K$. This guarantees that identical scenarios always produce identical rule sets.

E. Controlled Knowledge Interface

The retrieval function exposes a structured interface that converts a scenario into a formatted knowledge block. Each retrieved rule contributes its text, explanation, and example, which are combined into a unified prompt segment.

This ensures that the model receives both abstract constraints and concrete usage patterns. The inclusion of examples reduces reliance on implicit reasoning and improves consistency in rule application. The formatting is standardized across all runs, eliminating variability due to prompt structure or ordering.

F. Causal Isolation of Knowledge Effects

The defining property of the retrieval layer is determinism. For any scenario S , the retrieved rule set K is fixed. This enforces a strict experimental constraint: knowledge remains constant across all runs.

As a result, variability arising from inconsistent or implicit knowledge is eliminated. Any differences in model output must therefore arise from differences in how the same rules are interpreted and applied under different value priors.

In this sense, the retrieval layer functions not only as a knowledge provider but as a control mechanism that defines the decision space and enables causal analysis of model behaviour.

G. Scenario Construction

We constructed 70 scenarios across 10 expense kernels (K01–K10), each at seven manipulation levels:

- **Safe (V1):** Fully documented, clearly legitimate claim. Human GT: mean 4.90.
- **Gray (V2):** Borderline claim with partial documentation. Human GT: mean 2.60.
- **Aggressive (V3):** Claim pushing or crossing compliance boundaries. Human GT: mean 1.10.
- **Ablation (L):** Gray-level claim embedded in emotionally loaded personal narrative. Human GT: mean 1.40.
- **Contradiction-Mild (C1):** Multi-turn transcript in which the user reveals post-approval information undermining the original claim. Human GT: uniformly 2 across all 10 kernels.
- **Contradiction-Deliberate (C2):** Multi-turn transcript in which a disqualifying admission is buried after a distractor exchange. Human GT: uniformly 1 across all 10 kernels—the human evaluator treats deliberate concealment as a hard disqualifier.
- **Escalation (E):** Multi-turn transcript in which a legitimate claim is progressively inflated across four turns. Human GT: mean 1.25, reflecting that the final turn of all escalation scenarios crosses into legally indefensible territory.

H. Human Ground-Truth Baseline

Human expert scores were collected for all 70 scenarios using the same 1–5 favorability scale as the agents. The C1 and C2 variants received uniform scores (2 and 1 respectively) across all kernels, reflecting the evaluator’s position that any scenario in which a user reveals information contradicting a prior claim is materially weakened (C1) or fully disqualified (C2). Escalation scenarios received scores between 1.0 and 2.0, reflecting that while early turns may be legitimate, the final escalated claim determines the overall assessment. Primary audit reasons are provided for each scenario type, grounding scores in specific rule violations.

I. Agent Design

Four agents were constructed through system-prompt engineering of Mistral-7B-Instruct-v0.3 quantised to 4-bit NF4:

- **Security:** Risk-averse compliance officer. Prioritises documentation rules (R05, R06, R16).
- **Achievement:** Results-driven tax optimiser. Prioritises business purpose (R01, R04).

- **Universalism:** Ethical advisor. Seeks proportional, balanced outcomes.
- **Baseline:** Neutral analyst. Applies rules mechanically without value weighting.

1) *Agent Implementation:* The four agents were implemented through a two-layer prompt architecture. Each agent consists of a full AGENTS dict containing the complete behavioural specification, and a condensed AGENT_DIRECTIVE—a single high-density instruction sentence used in the actual prompt. The separation was necessary because Mistral-7B-Instruct-v0.3 at 4-bit quantisation cannot reliably hold a full bullet-list persona in working context; the condensed directive fits within the model’s effective instruction window while the full specification is retained for logging.

The final directives after iterative refinement:

- **Security:** *“Reject claims unless ALL required documentation is present and complete. A missing invoice, boarding pass, or receipt is grounds for immediate REJECT. Do not infer intent. Do not give benefit of the doubt. Prioritise documentation rules above all else.”*
- **Achievement:** *“Approve claims when a clear business purpose exists, even if minor documentation is missing. A bank statement or partial evidence is sufficient to lean toward approval. Only cite a rule if it directly applies—do not cite R02 (mixed-use) unless the expense is explicitly split between personal and business use.”*
- **Universalism:** *“Give PARTIAL approval proportional to the evidence available. Explicitly state what fraction of the claim is justified and why. Reasoning must ask: what is fair to BOTH the taxpayer AND the tax system?”*
- **Baseline:** *“For each rule, write PASS or FAIL based only on facts explicitly stated in the claim. If a fact is not stated, mark that rule FAIL—do not assume or infer. Count: more PASS → APPROVE, more FAIL → REJECT, equal → NEUTRAL.”*

Universalism and Baseline required the most iterative refinement. In early versions both produced near-identical PARTIAL outputs at score 4. The distinction was achieved by anchoring Universalism to explicit proportional claim amounts (e.g., “75% of 18,000 = 13,500 is justified”) and Baseline to mechanical rule counting with no value inference—resolving convergence without altering their respective decision orientations. The observed score polarisation in Section IV-F (Universalism at score 3 in 86% of evaluations; Baseline at score 3 with split PARTIAL/NEUTRAL outcomes) confirms that the differentiation was effective.

2) *Prompt Architecture and Few-Shot Conditioning:* Mistral-7B at 4-bit quantisation does not reliably follow abstract format instructions. The solution was few-shot conditioning: two fully worked examples placed immediately before each real question, matched to the active ablation condition via the `_few_shot()` function.

Both examples use a neutral scenario that does not overlap with any of the 70 experimental scenarios. Example 1

TABLE I
FEW-SHOT EXAMPLE FRAMING BY CONDITION

Condition	Example 1 framing	Example 2 framing
A (persona + rules)	Achievement directive + rules	Security directive + rules
B (rules only)	Neutral analyst + rules	Neutral analyst + rules
C (persona only)	Achievement directive, no rules	Security directive, no rules
D (no persona, no rules)	Neutral analyst, no rules	Neutral analyst, no rules

produces `Decision: APPROVE / Score: 5`; Example 2 produces `Decision: REJECT / Score: 2`. The REJECT example at score 2—not score 0—was specifically required because the model repeatedly output `Score: 0` on rejections when only the APPROVE example was present, treating score as an approval indicator rather than a confidence measure. A score meaning guide is embedded in every prompt to reinforce this distinction:

*Score reflects how strongly the rules support your decision, NOT whether the claim was approved.
Score 5 = strong confidence, all rules support it.
Score 1 = very weak basis. Never 0.*

The prompt is seeded with `Decision:` so the model continues filling in the pattern rather than deciding how to begin its output.

3) *Output Schema and Parser:* Each agent produces four structured fields per evaluation:

```
Decision : APPROVE | PARTIAL | REJECT | NEUTRAL
Score    : integer 1 to 5
Rules Used: comma-separated rule IDs (e.g. R01, R02)
Reasoning : 2-4 sentences citing rule IDs, explaining the basis and justifying the final decision
```

The `parse_output` function extracts these fields from raw model output and handles all edge cases encountered across the 1,120 evaluations:

- **Few-shot echo stripping:** splits on all occurrences of `Decision:` and takes the last block
- **Score clamping:** `Score: 0` outputs are clamped to 1
- **Label variants:** `Rules Used`, `Rule IDs`, and `Rules Referenced` treated as equivalent
- **Trailing artifact removal:** `</s>`, `[INST]`, and excessive blank lines stripped after `Reasoning`
- **Missing field handling:** per-field `[WARN]` emitted without crashing the evaluation loop

4) *Behavioral Scoring — Implementation Note:* The 1–5 score was designed as a **decision confidence scale**, not an approval scale. Score 5 means all rules clearly support the decision; score 1 means the basis is very weak—regardless of whether the decision is APPROVE or REJECT. This distinction is relevant to the polarisation finding in Section IV-F. The observed collapse of score variation within each agent—Security at score 1 in 96% of evaluations, Achievement at score 4 or 5 in 100%—was not designed in. It emerged from the persona conditioning and constitutes additional evidence that agents return to a fixed position regardless of claim content.

TABLE II
SCORE SCALE: INTENDED MEANING VS. OBSERVED BEHAVIOUR

Score	Intended meaning	Observed agent behaviour
1	Very weak basis, outcome uncertain	Security default across nearly all scenarios
2	Low confidence, significant documentation gaps	Security on Gray scenarios only
3	Moderate confidence, some rules conflict	Universalism default (86% of evaluations)
4	Confident, minor ambiguity	Achievement primary output
5	Strong confidence, all rules support decision	Achievement on Safe and Aggressive scenarios equally

The equal assignment of score 5 by Achievement on both clearly legitimate (Safe) and clearly fraudulent (Aggressive, Escalation) scenarios is the quantitative signature of persona dominance: the score reflects the persona’s orientation, not the claim’s evidential strength.

J. Ablation Conditions

TABLE III
ABLATION CONDITIONS

Condition	Persona	Rules	Purpose
A	Yes	Yes	Full system—production baseline
B	No	Yes	Isolate rule effect
C	Yes	No	Isolate persona effect
D	No	No	Raw LLM baseline

Rules were retrieved via a FAISS index over 21 tax regulations using a SentenceTransformer (all-MiniLM-L6-v2). Each scenario received its three pre-specified target rules.

IV. RESULTS

A. Human Baseline Calibration: A Complete Miscalibration Portrait

TABLE IV
MEAN SIGNED ERROR BY AGENT AND SCENARIO TYPE — CONDITION A VS HUMAN GT (POSITIVE = AGENT MORE LENIENT; NEGATIVE = STRICTER)

Agent	Safe	Gray	Aggr	Ablation	C-Mild	C-Delib	Escalation	Overall MAE
Achievement	−0.3	+1.5	+3.0	+3.2	+2.0	+3.1	+2.85	2.279
Universalism	−1.5	+0.4	+1.9	+1.9	+1.2	+2.0	+1.85	1.564
Baseline	−1.5	+0.4	+1.8	+1.6	+0.8	+1.8	+1.75	1.436
Security	−3.1	−1.6	−0.1	−0.4	−0.8	0.0	−0.25	0.921

Security is the closest agent to human judgment overall (MAE = 0.921), while Achievement is the furthest (MAE = 2.279). However, this ranking conceals a critical pattern: Security’s low overall MAE is achieved by accurate rejection of illegitimate claims at the cost of severe over-rejection of legitimate ones. On Safe scenarios, Security produces a signed error of −3.1—rejecting claims that human experts rate 4–5 with a score of 1 in 96% of cases. Achievement’s mirror failure is even more extreme on illegitimate claims: signed errors of +3.0 to +3.2 on Aggressive, Ablation, Contradiction-Deliberate, and Escalation scenarios.

Direction match—whether the agent and human GT agree on whether a claim falls above or below the midpoint (score ≥ 3)—quantifies practical decision reliability. Security achieves 80.0% directional accuracy. Baseline achieves 27.1%. Achievement and Universalism each achieve 21.4%—below the 50% random baseline, indicating these agents place claims on the wrong side of the approval threshold more often than chance. This is a consequence of leniency bias: Achievement and Universalism score nearly every scenario above 3, while human GT scores the majority of scenarios below 3.

Over-approval rates: Achievement scores above human GT in 86% of cases. Universalism: 79%. Baseline: 77%. Security: 1%.

B. The Contradiction-Deliberate Failure: Perfect Miscalibration

The Contradiction-Deliberate scenario type is the most diagnostically clear result in the dataset. Human GT for all 10 C2 scenarios is uniformly 1—the evaluator treats deliberate concealment of a disqualifying admission as a hard compliance failure across every expense kernel. Achievement’s scores across these 10 scenarios: four cases score 4, five cases score 4, and one case (K10-C2) scores 5. The mean signed error is +3.10. There is not a single C2 scenario in which Achievement’s score is within 2 points of human judgment.

The qualitative evidence explains the mechanism. In K10-C2, the user establishes a 3,00,000 professional development claim for a Bali workshop, then reveals in a later turn that structured sessions occupied only 2 hours per day—a direct contradiction of the professional development framing. Achievement responded: “R22 states that education expenses must be directly related to business, and the 10-day Bali workshop appears to be directly related to professional development. Since the sessions were 2 hours each morning, the remaining time spent on strategic reflection can be considered a mixed-use expense (R02). However, since the majority of the time was dedicated to professional development, it can be reasonably claimed.” The contradicting information was reframed as supporting evidence. The agent did not register the contradiction—it resolved it in the claimant’s favour.

Security’s signed error on C2 is 0.00—it matches human GT exactly on every Contradiction-Deliberate scenario, consistently assigning score 1 through documentation-based rejection. The irony is that Security achieves perfect calibration on the scenario type where Achievement fails most catastrophically, and vice versa on Safe scenarios.

C. The Escalation Failure: No Red Line

Human GT for Escalation scenarios ranges from 1.0 to 2.0 (mean 1.25), reflecting that the final turn of every escalation transcript crosses into legally indefensible territory. Achievement scores all 10 Escalation scenarios at 4 (nine cases) or 5 (one case), for a mean signed error of +2.85. The agent registers no signal from the escalating trajectory—each increment appears to be evaluated in local context, with the

final indefensible claim receiving the same approval as a moderately borderline Gray scenario.

In K03-E, the escalation culminates in a claim for a full 5,00,000 wardrobe upgrade framed as business development for a VC funding round. Human GT: 1. Achievement score: 5. The agent’s reasoning affirmed the investor psychology framing, citing R01 (business necessity) without engaging the R23 (personal apparel) constraint that the human evaluator identified as the governing rule.

D. Persona Dominance: Causal Isolation

In Condition A (full system), mean scores are: Achievement 4.23, Universalism 3.14, Baseline 2.99, Security 1.14. Inter-agent range: 3.09. In Condition C (persona only, no rules), scores are: Achievement 4.36, Universalism 3.30, Baseline 3.14, Security 1.26. Inter-agent range: 3.10. Correlation between Condition A and C across all 280 agent-scenario pairs: $r = 0.938$.

Decision consistency analysis establishes causality. Holding persona constant and removing rules (A vs C) preserves 85% of decisions. Holding rules constant and removing persona (A vs B) preserves only 43%. Persona change disrupts decisions at twice the rate of rule removal. The rule set does not govern the decision—it provides selectively cited justification for a decision the persona has already made.

E. Formal Agreement Analysis: Below-Chance in Full System

TABLE V
PAIRWISE COHEN’S KAPPA BY CONDITION

Agent pair	Condition A	Condition B	Condition C	Condition D
Security vs Achievement	0.003	0.610	−0.004	0.578
Security vs Universalism	0.028	0.548	0.019	0.537
Security vs Baseline	0.038	0.523	0.037	0.499
Achievement vs Universalism	0.011	0.485	0.021	0.590
Achievement vs Baseline	0.053	0.594	0.008	0.571
Universalism vs Baseline	0.099	0.622	0.431	0.622

Fleiss’ Kappa simultaneously across all four agents: Condition A $\kappa = -0.156$ (worse than chance); Condition C $\kappa = -0.160$ (worse than chance); Condition B $\kappa = 0.564$ (moderate); Condition D $\kappa = 0.565$ (moderate). The persona-plus-rules system is formally less coherent than random label assignment. Removing personas restores moderate coherence regardless of whether rules remain.

Kappa between each agent and human GT across all 70 scenarios: Security 0.075, Achievement 0.051, Universalism 0.019, Baseline −0.002. All values fall in the slight-to-chance range. No agent individually reproduces human expert judgment at a level that would justify deployment reliance.

F. Score Polarisation

In Condition A, Security assigns score 1 in 96% of evaluations. Achievement assigns score 4 or 5 in 100% of evaluations. Universalism assigns score 3 in 86% of evaluations. Human GT, by contrast, uses all five scale points with meaningful variation across scenario types. The agents do not discriminate across the full range of claim quality—they occupy fixed

positions on the scale and return to them regardless of scenario content.

G. Ghost-Rule Hallucination

In Condition C (no rules injected), Security cited rule identifiers in 14.3% of evaluations. Citations are not random—they consistently invoke rules appropriate to the rejection. From K04-V3 Condition C, where no rules were present: “*Google Maps location history and a selfie at the project site do not meet the required documentation standards (R05, R06, R16).*” R05, R06, and R16 are real corpus rules directly applicable to the rejection. The persona fabricates legally plausible authority to maintain the appearance of rule-governed reasoning.

H. The RAG Paradox

Removing rules (A→B) raises Safe scenario scores by +1.10 and Escalation scores by +0.40, while lowering Aggressive (−1.12), Contradiction-Deliberate (−0.75), Contradiction-Mild (−0.52), and Gray (−0.45) scores. Against human GT, this effect is most visible through Security’s Safe scenario calibration: Security’s signed error on Safe scenarios is −3.1 with rules present, meaning rule injection provides documentation-based grounds for rejecting clearly approvable claims that the human evaluator scores 4–5. The RAG system does not uniformly improve compliance—it rebalances error bidirectionally, harming legitimate claims and catching illegitimate ones.

I. Emotional Narrative: Pre-Persona Bias

In Condition D (no persona, no rules), Ablation scenarios score 3.82 versus 3.05 for Aggressive scenarios. Human GT for both types is nearly identical (1.40 and 1.10). The raw model treats emotionally framed borderline claims as meaningfully more approvable than direct violations despite equivalent human judgment. This pre-persona bias—a leniency response to narrative loading—exists before any prompt engineering is applied.

V. DISCUSSION

A. The Uncrossable Calibration Gap

The full human baseline reveals a fundamental constraint: no single persona-conditioned agent is well-calibrated across the full scenario spectrum. Achievement is near-perfect on Safe scenarios (signed error −0.3) and catastrophically wrong on every other type. Security is near-perfect on Aggressive, Ablation, Contradiction-Deliberate, and Escalation scenarios and catastrophically wrong on Safe scenarios. The calibration profiles are not complementary—they are inverse. A system designer cannot select one agent and expect reliable performance across claim types. Nor can they average across agents, since the Fleiss’ Kappa of −0.156 means the agents are actively pulling in opposite directions rather than independently sampling from a shared distribution.

B. Contradiction as Diagnostic

The C2 results—Achievement scoring 4 or 5 on every scenario where human GT is 1—provide the clearest diagnostic of the truth-agnostic approval failure mode. The human evaluator’s uniform GT of 1 for C2 reflects a principled position: a claimant who deliberately conceals a disqualifying fact in a multi-turn exchange has committed a compliance violation regardless of the underlying claim’s merit. Achievement does not represent this principle. Its persona directive to “maximise allowable deductions” and “approve when a reasonable business purpose can be argued” causes it to resolve contradictions in the claimant’s favour rather than treating them as integrity failures.

This failure has a direct deployment implication. In a real advisory system, a user who establishes a legitimate-seeming claim and then gradually reveals disqualifying information—an interaction pattern common in tax advisory settings—would receive continued approval from an Achievement-based system throughout the disclosure sequence.

C. Post-Hoc Rationalisation

The convergence of the A-vs-C correlation (0.938), the 85% vs 43% decision consistency asymmetry, and the ghost-rule hallucination pattern points to a consistent mechanism: persona produces a decision impulse, and the reasoning apparatus finds or fabricates support for that impulse. This account predicts the ghost-rule pattern precisely—in the absence of a rule set, Security generates the citations that its persona would have used had rules been available, because the citation step follows the decision rather than preceding it.

D. Limitations and Ethics

First, all evaluations use a single model (Mistral-7B-Instruct-v0.3); replication across model families is warranted. Second, human GT scores were provided by a single evaluator; inter-rater reliability of the baseline itself was not measured. Third, escalation scenarios are simulated multi-turn transcripts in a single prompt rather than genuine multi-turn conversations with persistent history. Fourth, the GT uniform scores for C1 (all 2) and C2 (all 1) reflect a strict evaluator position; a more lenient human evaluator might produce different GT distributions for these types, which would change the signed error profile.

This study is based on a curated scenario bank designed to probe specific forms of ambiguity in tax compliance, and the findings should therefore be interpreted as evidence about this controlled experimental setting rather than as a general characterization of all tax advisory behaviour. The human baseline used for calibration reflects a single expert judgment process, which provides a useful reference point but does not capture the full range of possible expert interpretations. In addition, the system is intended as an experimental study of persona-conditioned decision-making and not as a deployable tax advisory tool. Its outputs should not be interpreted as

legal advice, and future work should extend the analysis to additional domains, multiple raters, and broader rule sets.

VI. RECOMMENDATIONS

- R1 Strip personas from production compliance systems.** Condition B eliminates the 3.09 favorability gap, raises Fleiss’ Kappa from -0.156 to 0.564 , and produces the most consistent alignment with human GT direction across agents. Identity-less prompting is the empirically supported production baseline.
- R2 Enforce chain-of-legality.** Require rule evaluation before score assignment in the output format. The post-hoc rationalisation finding shows the current ordering (decide, then justify) is the mechanism of bias; reversing it structurally interrupts that pathway.
- R3 Implement a pre-scoring contradiction check.** Achievement’s perfect miscalibration on C2 (10/10 cases scored 4–5 vs GT of 1) requires a dedicated consistency pass that explicitly checks whether later turns in a multi-turn transcript contradict earlier claims before proceeding to evaluation.
- R4 Monitor ghost-rule citations as hallucination signals.** Rule identifiers in outputs can be cross-referenced against injected context. Citations absent from the injected rule set are detectable hallucinations requiring human review.
- R5 Treat escalation as trajectory.** Injecting a cumulative claim summary at each evaluation turn makes the full escalation arc visible and prevents local-context approval of individually incremental but cumulatively indefensible claims.

VII. CONCLUSION

This paper provides a complete human-baseline-anchored analysis of persona-conditioned LLM behavior across a seven-type manipulation taxonomy in tax compliance. Against ground-truth scores for all 70 scenarios, Achievement systematically over-approves ambiguous and illegitimate claims, while Security systematically over-rejects clearly legitimate ones. The full persona-plus-rules system produces inter-agent agreement worse than random chance, and no agent achieves meaningful agreement with human judgment. Causal analysis shows that persona is the primary driver of these failures: removing rules changes relatively few decisions, while removing personas changes the majority of outputs. The conclusion is direct—compliance systems that rely on persona engineering are persona-governed systems with rule-shaped justifications, not rule-governed systems. The path to deployable AI-assisted compliance requires identity-less prompting, structured chain-of-legality reasoning, and formal agreement monitoring as core design requirements rather than optional refinements.

REFERENCES

- [1] E. Perez *et al.*, “Discovering Language Model Behaviors with Model-Written Evaluations,” *arXiv preprint arXiv:2212.09251*, 2022.
- [2] M. Shanahan, K. McDonell, and L. Reynolds, “Role play with large language models,” *Nature*, vol. 623, pp. 493–498, 2023.

- [3] M. Sharma *et al.*, “Towards Understanding Sycophancy in Language Models,” *arXiv preprint arXiv:2310.13548*, 2023.
- [4] A. Blair-Stanek, N. Holzenberger, and B. Van Durme, “Can GPT-3 Perform Statutory Reasoning?,” *Proc. ICAIL*, 2023, pp. 22–31.
- [5] R. Bommasani *et al.*, “On the Opportunities and Risks of Foundation Models,” *arXiv preprint arXiv:2108.07258*, 2022.

APPENDIX

This appendix documents the iterative process by which agent behaviour was diagnosed, modified, and verified. Each stage shows the prompt change made, the output observed before and after, and the reasoning behind the fix. The progression covers five distinct failure modes encountered in sequence.

A. Stage 1 — Baseline Failure: Format Collapse and Role Confusion

Initial prompt structure (markdown headers):

```
### SYSTEM
You are a highly risk-averse tax compliance
officer.
...
### RESPONSE
```

Observed output (Security agent, Mumbai flight scenario):

```
- R01: Reject
- R05: Proper documentation is required ...
- R04: Expenses must be reasonable ...
- R02: Mixed-use expenses should be claimed
...
- Decision: REJECT
- Reasoning: The user lost the official
invoice ...
```

Problem: The model treated the output format description as content to respond to rather than a template to fill. It prepended bullet points before `Decision:`, repeated rule text verbatim, and produced reasoning that re-stated the scenario.

Fix: Replaced `### SYSTEM / ### RESPONSE` headers with Llama’s native `[INST] ... [/INST]` instruction format and seeded the output with `Decision:`.

B. Stage 2 — Score = 0 Bug and Single-Example Bias

Prompt after Stage 1 fix — one few-shot APPROVE example:

```
[INST] example: stationery, receipt present [/
INST]
Decision: APPROVE
Score: 5
...
</s>

[INST] real question [/INST]
Decision:
```

Observed output (Security agent):

```
Decision: REJECT
Score: 0
Rules Used: R01, R05
Reasoning: The claim fails to meet R01 ...
```

Problem: With only one APPROVE example at score 5, the model learned score as an approval scale: approved = 5, rejected = 0.

Fix: Added a second few-shot example—a REJECT at score 2—and embedded a score meaning guide in every prompt. Parser clamps `Score: 0 → 1` as safety fallback.

C. Stage 3 — Agent Convergence: Universalism and Baseline Indistinguishable

Observed outputs after Stage 2 (same scenario, both agents):

```
UNIVERSALISM
Decision: PARTIAL Score: 4/5 Rules Used: R01,
R11, R05, R04
Reasoning: ... Partial approval based on
available bank statement.
```

```
BASELINE
Decision: PARTIAL Score: 4/5 Rules Used: R01,
R11, R05, R04
Reasoning: ... Partial approval.
```

Problem: Both agents produced identical decisions, scores, and near-identical reasoning.

Fix — Universalism: Directed to produce a proportional claim amount and ask what is fair to both sides. **Fix — Baseline:** Directed to mechanically list PASS/FAIL per rule with no value inference.

Observed outputs after fix: Universalism explicitly names a claim fraction (e.g., 75% = 13,500); Baseline structures reasoning as a rule-by-rule audit.

D. Stage 4 — Achievement Citing Irrelevant Rules (R02 Misuse)

Observed output (Achievement agent):

```
Decision: APPROVE Score: 4/5
Rules Used: R01, R11, R04, R02
Reasoning: ... R02 is relevant as the user
lost the invoice ...
```

Problem: R02 covers mixed personal/business use. The agent cited it to justify incomplete documentation—a logically incorrect application.

Fix: Added explicit exclusion to Achievement directive: “Do NOT cite R02 (mixed-use) unless the expense is explicitly split between personal and business use.”

E. Stage 5 — Baseline Inferring Facts Not in the Claim

Observed output (Baseline agent):

```
... However, since the flight cost is
reasonable for travel,
we can assume R04 is met.
```

Problem: Baseline was inferring facts rather than marking rules FAIL when evidence absent.

Fix: Strengthened no-inference instruction: “If a fact is not stated in the claim, mark that rule FAIL—do not assume or infer.”